

# FLT Seminar Series<sup>1</sup>, Session 2

## How Feature Learning Theory Works?

Chen Yanbo<sup>2</sup>

Ph.D. Candidate  
School of Computer Science, Wuhan University  
430072, Wuhan, Hubei, China

Jun. 1st, 2025



---

<sup>1</sup>This project is open for collaboration. For details, see our project page at <https://github.com/yanboc/feature-learning-theory>.

<sup>2</sup>Contact: [yanboch@126.com](mailto:yanboch@126.com).

## An Quick Start to FLT

- Highlights from our last session: **what** is feature learning theory?
  - ▶ Terminologies: what are feature and learning, respectively?
  - ▶ A bird's-eye view summary of FLT
- A simplified example: **how** FLT works?
  - ▶ The **theoretical framework** of FLT
  - ▶ The **theoretical goal** of FLT

# Table of Contents

1 A brief review of our last session


2 A simplified example: how FLT works?

# What is *feature learning theory (FLT)*?

What are *feature* and *learning*, respectively?

## Terminologies

- We focus on **features** in DL (w.r.t. data, NNs, and specific tasks); the main goal of DL is to *find NNs that extract useful features from data*.



**Figure:** Higher- and lower-level features in CNN-based classification

# What is *feature learning theory (FLT)*?

What are *feature* and *learning*, respectively?

## Terminologies

- Machine **learning** uses a specified *algorithm* to find the best model in the *hypothesis class* (e.g., NNs) according to the performance of the model on the *data*, concerning the *evaluation* standard.

**Table:** The *four core elements* of the ML&DL paradigm

	<b>Theoretical</b>	<b>In Practice</b>
<b>Data</b>	vectors and matrices	tensor
<b>Hypothesis Class</b>	functions and mappings	multi-layer NN
<b>Algorithm</b>	optimization	optimizer, LR, ...
<b>Evaluation</b>	loss function, regularization	CE, MSE, ...

# A bird's-eye view summary of FLT

**Machine learning** uses a specified *algorithm* to find the best model in the *hypothesis class* (e.g., NNs) according to the performance of the model on the *data*, concerning the *evaluation* standard.



**FLT** specifies the learning task (*network structure*, *data assumption*, *loss*, and *algorithm*) and explore the **dynamics** of training.



**Dynamics**: how the *parameters of the NN* iterate from random initialization (noise) to *useful features* capable of accurate classification/regression?

# Table of Contents

- 1 A brief review of our last session
- 2 A simplified example: how FLT works?

# How FLT works? (1/2)

Step 1. FLT *specifies* the learning task  
(*network structure, data assumption, loss, and algorithm*)

## Theoretical framework [Allen-Zhu & Li, 2005.10190v4]

- **Hypothesis Class:** 2-layer (symmetric)-ReLU network  $f(x; w)$
- **Data:** orthogonal feature + sparse coding model

$$x = Mz + \xi, \quad y = \text{sign}(\langle w^*, z \rangle)$$

- **Algorithm:** GD with random initialization

$$w^{(t+1)} = w - \eta \nabla \text{Loss}(f(x, w), y)$$


- **Evaluation:** logistic loss for classification.

**Intuition:** specifying the learning task is like creating a **virtual environment** to *play around* with.



# The triviality & tractability trade-off

Specifying the learning task is a **tricky job**



**Figure:** There is a trade-off between triviality and tractability.

# Data Distribution

We consider a supervised binary classification task.

## Sparse Coding Model, SCM (1/2)

Consider the training data  $x \in \mathbb{R}^d$  generated from  $x = \mathbf{M}z + \xi$  for *coding matrix*  $\mathbf{M} \in \mathbb{R}^{d \times d}$ , *hidden vector*  $z \in \mathbb{R}^d$ , and *noise vector*  $\xi \in \mathbb{R}^d$  such that

- $\mathbf{M} = (\mathbf{M}_1, \mathbf{M}_2, \dots, \mathbf{M}_d)$  is a unitary (i.e., orthonormal) matrix.
- $z$  is a sparse vector in the sense that  $\|z\|_0 = \Theta(\sqrt{d})$ .
- for simplicity, we let  $\xi \sim \mathcal{N}(0, \sigma_0 \mathbf{I})$  in this seminar

### Remarks:

- The *sparsity* of  $z$  and the *magnitude* of  $\xi$  are *non-negligible* to guarantee the expressiveness of SCM.
- More choices of (implicit) data distributions (delayed to FLT-3).

## Sparse Coding Model, SCM (2/2)

The label of  $x$  is decided by the hidden vector  $z$  and a *labeling function*  $w^\star$ .

$$y = \text{sign}(\langle w^\star, z \rangle). \quad (1)$$

For simplicity, assume that  $|w_i^\star| = \Theta(1)$  for all  $i \in [d]$  (i.e., *balanced* setting).

**learning goal:** predict the label of  $x$



We need to find  $f$  such that  $f(Mz + \xi) \approx \text{sign}(\langle w^\star, z \rangle)$ .

# The Philosophy of FLT

## Philosophy No.1., **Symmetry**

In the specified learning tasks, *the data, networks, and algorithms must embody certain symmetries or self-similarities*. For instance,

- the coding matrix  $\mathbf{M}$  is orthonormal,
- the labeling function  $w^\star$  is balanced,
- the self-similarity within the GD algorithm.

In FLT, we only need to analyze a single part of the symmetric system, rather than all the parts.

# Hypothesis Class (i.e., Network Structure)

**learning goal:** predict the label of  $x$



We need to find  $f$  such that  $f(\mathbf{M}z + \xi) \approx \text{sign}(\langle \mathbf{w}^*, z \rangle)$ .

## Natural intuition

A natural intuition is letting

$$f(x) = \langle \mathbf{w}^*, \mathbf{M}^\top x \rangle = \langle \mathbf{w}^*, z \rangle + \langle \mathbf{M}\mathbf{w}^*, \xi \rangle. \quad (2)$$

The non-negligible magnitude of  $\xi$  would lead to inaccuracy.



Linear models are *not complicated (expressive) enough* to characterize SCM.

(cf. the expressiveness of NN, VC-dimension theory, the overfitting phenomenon, regularization & the Occam's Razor principle.)

# Hypothesis Class (i.e., Network Structure)

## Two-layer (symmetric) ReLU network

We consider the following network

$$f_t(x; \mathbf{w}^{(t)}) = \sum_{i=1}^m \left( \text{ReLU} \left( \langle \mathbf{w}_i^{(t)}, x \rangle - b_i^{(t)} \right) - \text{ReLU} \left( -\langle \mathbf{w}_i^{(t)}, x \rangle - b_i^{(t)} \right) \right)$$

(optimized to)  $f(x) \approx \sum_{i=1}^n w_i^* \left( \text{ReLU}(\langle \mathbf{M}_i, x \rangle - b_i) - \text{ReLU}(-\langle \mathbf{M}_i, x \rangle - b_i) \right)$

parameterized by  $\mathbf{w}^{(t)} := \left( w_{[m]}^{(t)}, b_{[m]}^{(t)} \right)$ , where  $m$  denotes the width of  $f_t$ .

### Remarks:

- The ReLU activation is smoothified (using a mollifier), omitted here.
- **How to obtain Eq. (14)?** It can neither be more complicated nor simpler (cf. Figure 2). *Pure tricks or intuition, maybe.*
- Over-parameterization & Thresholding.

# The Philosophy of FLT

- Philosophy No.1., Symmetry.

## Philosophy No.2., **Programmatic Thinking**

When performing feature learning analysis, *one should think and act like a programmer, rather than a mathematician.* For instance,

- **Programmatic definitions.** Find intuitions and definitions from practical code and PyTorch documentation!
- **Programmatic tuning.** There are many parameters in the analysis, e.g.,  $m$  and  $\sigma_0$ , that require careful tuning.
- **Programmatic workflow.** FLT undergoes the entire training process, starting with random initialization and stopping by the attainment of an accurate classifier.

In FLT, we only commit the *minimum necessary changes* to a practical training process of NNs.

# Evaluation

## Loss function and empirical risk

We consider the standard *logistic loss*

$$\text{Loss}_t(\mathbf{w}^{(t)}; x, y) := \log \left( 1 + \exp \left( -y f_t(x; \mathbf{w}^{(t)}) \right) \right) \quad (3)$$

and the corresponding *empirical risk*

$$\text{Risk}_t(\mathbf{w}^{(t)}) := \frac{1}{N} \sum_{j \in [N]} \left( \text{Loss}_t(\mathbf{w}^{(t)}; x_j, y_j) \right) \quad (4)$$

The **training goal** is to find the best  $\mathbf{w}$  that minimizes the empirical risk (i.e., the ERM training paradigm).

### Remark:

- FLT for other training paradigms (e.g., Bayesian NN, GANs, RL, Causal Inference). (Good choices for future research!)



# Algorithm

## Gradient descent with random initialization

For each  $i \in [m]$ , the update rule of  $\mathbf{w}_i^{(t)}$  is

$$\mathbf{w}_i^{(t+1)} \leftarrow \mathbf{w}_i^{(t)} \eta \nabla_{\mathbf{w}_i^{(t)}} \text{Risk}(\mathbf{w}_i^{(t)}), \quad (5)$$

for any  $t \in [T]$ , where  $\mathbf{w}$  is initialized as

$$\mathbf{w}_i^{(0)} \sim \mathcal{N}(0, \sigma_1^2 \mathbf{I}). \quad (6)$$

### Remark:

- About the parameters: recall the *programmatic thinking* philosophy.
- Some neurons **have already been good enough** at initialization (cf. *concentration inequalities* & the *lottery ticket hypothesis*).

# Summary

- The first step of FLT is to *specify* the learning task, including network structure, data assumption, loss, and algorithm (check it!).
- Specifying the learning task is like creating a **virtual environment** to *play around* with.
  - ▶ What is *playing around*? Acts like tuning parameters, network structures, and data assumptions.
  - ▶ How to advance an FLT proof? Just play around and observe the changes in the proofs. The *difficulty curve* of FLT is almost linear.
- The Philosophy of FLT No. 1&2
  - ▶ Design a symmetric system to reduce the complexity of analysis.
  - ▶ Think and act like a programmer, rather than a mathematician.

## How FLT works? (2/2)

Step 2. FLT defines multiple **good property sets** and studies how the neurons **enter or exit** these sets. (**Dynamics!**)

### What *defines* a good feature?

Recall the network structure

$$f_t(x; \mathbf{w}^{(t)}) = \sum_{i=1}^m \left( \text{ReLU} \left( \langle \mathbf{w}_i^{(t)}, x \rangle - b_i^{(t)} \right) - \text{ReLU} \left( -\langle \mathbf{w}_i^{(t)}, x \rangle - b_i^{(t)} \right) \right)$$

(optimized to) 
$$f(x) \approx \sum_{i=1}^n \mathbf{w}_i^* \left( \text{ReLU} \left( \langle \mathbf{M}_i, x \rangle - b_i \right) - \text{ReLU} \left( -\langle \mathbf{M}_i, x \rangle - b_i \right) \right)$$

- One good feature  $\mathbf{w}_j^{(t)}$  should approximate the *direction* of  $\mathbf{M}_i$ .
- Multiple good features should approximate the *magnitude* of  $\mathbf{w}_i^*$ .

# Good Property Sets

FLT defines multiple levels of good property sets. We consider two of them.

## Surely Good Neurons $\mathcal{S}_{j,sure}^t$ and Potentially Good Neurons $\mathcal{S}_{j,pot}^t$

Let  $\mathcal{S}_{j,sure}^t \subseteq [m]$  be those neurons  $i \in [m]$  satisfying

- $\langle \mathbf{w}_i^{(t)}, \mathbf{M}_j \rangle^2 \geq (c_1 + c_2)(\sigma_3^{(t)})^2 \log d$ ,
- $\langle \mathbf{w}_i^{(t)}, \mathbf{M}_{j'} \rangle^2 < (c_1 - c_2)(\sigma_3^{(t)})^2 \log d$ , for every  $j' \neq j$ ,
- $\langle \mathbf{w}_i^{(t)}, \mathbf{M}_j \rangle \mathbf{w}_j^\star > 0$ .

Let  $\mathcal{S}_{j,pot}^t \subseteq [m]$  be those neurons  $i \in [m]$  satisfying

- $\langle \mathbf{w}_i^{(t)}, \mathbf{M}_j \rangle^2 \geq (c_1 - c_2)(\sigma_3^{(t)})^2 \log d$ .

### Remarks:

- For the flexibility of the theory, more parameters are introduced.
- **Feature Learning:** The neurons  $\{\mathbf{w}_i\}_{i \in \mathcal{S}_j}$  approximate the direction and magnitude of  $\mathbf{M}_j$  and  $\mathbf{w}_j^\star$ .

# How the neurons *enter and exit* these sets?

## Theoretical Goals

The *desired principles* of neurons' entering and exiting good property sets can be summarized as follows.

### Entering:

- Some of the neurons have already been in these sets at initialization.
- Neurons from lower-level sets enter higher-level sets with probability.

### Exiting:

- Neurons exit lower-level sets and enter higher-level sets.
- Neurons never exit the highest-level sets.

The **main goal** of FLT analyses are to prove the above principles.

The proof techniques are postponed to FLT-3.

# Thanks for your participation!



Welcome to join our WeChat group!

If this expires, please don't hesitate to contact me at [yanboch@126.com](mailto:yanboch@126.com).